



# Driverless Car Sentiment Analysis and Impact

Kyle Markwardt



# Objective

## Question:

What are the downstream impacts of paper-of-record headlines on the burgeoning autonomous vehicle industry?

Supposition is that the stock market reacts to headlines.

## Goals:

Conduct sentiment analysis on Autonomous Vehicle headlines

Create a model to predict impact on self-driving company stock prices based on volume and sentiment of headlines

Judge various company sensitivities to headlines about self driving



# Hurdles and Limitations

## Stock data:

- Many main players owned by much larger companies (Cruise:GM, Waymo:GOOG)
- Many failed or never went public
- Others are partial plays (TSLA, NVDA, APTV)
- Pure plays have short public history (AUR, TSPH)

## News:

- Tech Crunch has no API
- NewsAPI.org includes sources like TC and full text, but is limited to 1 month search
- NYT limited to headline and snippet



# Data Sources

{T} Developers

## Stock data

- Yfinance

## News

- NYT API
  - 500 articles for Autonomous Driving back to Jan 2016

## Labeled Training Data [Twitter self driving sentiment]

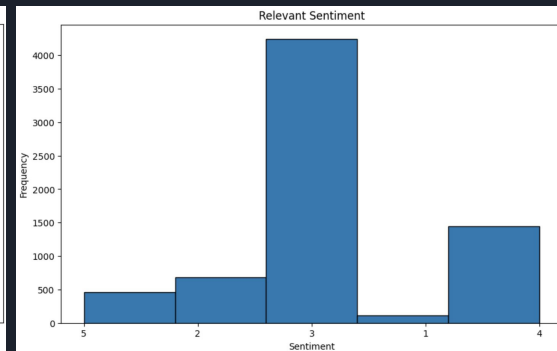
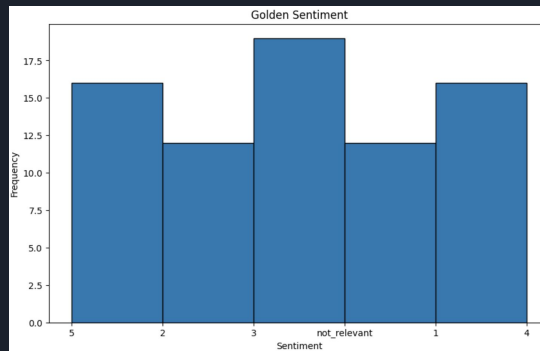
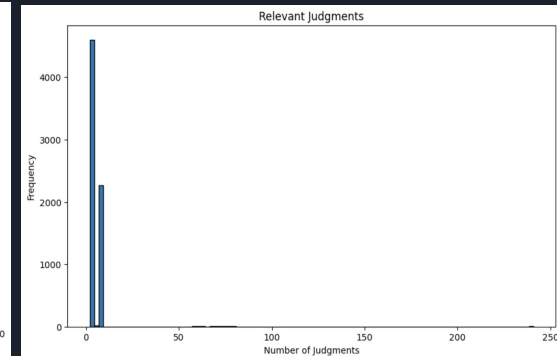
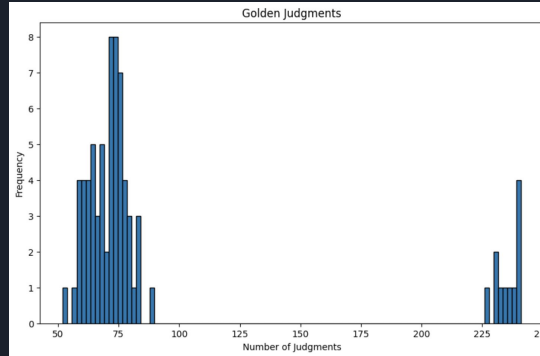
- <https://data.world/crowdfower/sentiment-self-driving-cars>
- 7000+ rows, created 2015

**yahoo!**  
finance

# Plan A - Create Sentiment Model

1. Create sentiment model on labeled twitter data
  - a. Convert 1-5 ratings to pos/neg
  - b. (drop 3s)
2. Predict headline sentiment with that model
  - a. Label smaller ( $n=20$ ) set of the headlines and evaluate model performance

Examining the twitter data:




# Result of Plan A:

Tokenized texts, max words 250, Embedding

- Dense Neural Network
- Bidirectional
- Dropout (overfitting was a huge problem with the small dataset)
- About 73%, not terrible...

Maximum length: 272  
Minimum length: 77



Epoch 10/10

54/54 [=====] - 65s 1s/step - loss: 0.0647 - accuracy: 0.9693 - val\_loss: 1.4321 - val\_accuracy: 0.7292

17/17 [=====] - 2s 134ms/step - loss: 1.4294 - accuracy: 0.7296

Test accuracy: 0.729629635810852

But when I ran the headlines evaluation set through it... 35%

1/1 [=====] - 0s 158ms/step - loss: 4.2869 - accuracy: 0.3500

Test accuracy: 0.3499999940395355



# Plan B: Pre-trained Sentiment Model

Decided to use a tried and trusted sentiment analysis model and work on the stock price correlation



```
sentiment_model = pipeline('sentiment-analysis', model='distilbert-base-uncased-finetuned-sst-2-english')
```

Voila!

Accuracy: 0.8

Classification Report:

	precision	recall	f1-score	support
-1	0.80	0.80	0.80	10
1	0.80	0.80	0.80	10
accuracy			0.80	20
macro avg	0.80	0.80	0.80	20
weighted avg	0.80	0.80	0.80	20



# Plan B continued

Features: (1) Sum of article sentiment aggregated by day (net sentiment per day)

- e.g.  $[(0.9564) \times (-1)] + [0.7854 \times (1)]$

Multiple models:


- Stock price on same day
- Stock price on next day
- Aggregate stock and sentiment by week
- Loop through multiple stocks

RNN with LSTM

Loss: MSE b/c of continuous nature of pct\_change



```
# Build and train LSTM model
def train_lstm(X_train, y_train, X_test, y_test):
    model = Sequential()
    model.add(LSTM(128, return_sequences=True, input_shape=(X_train.shape[1], 1)))
    model.add(LSTM(128))
    model.add(Dropout(0.2))
    model.add(Dense(1))
    model.compile(optimizer=Adam(learning_rate=0.001), loss='mean_squared_error')
    history = model.fit(X_train, y_train, batch_size=32, epochs=10, validation_data=(X_test, y_test), verbose=1)
    loss = model.evaluate(X_test, y_test)
    print(f'Test Loss: {loss}')
    y_pred = model.predict(X_test)
    mse = mean_squared_error(y_test, y_pred)
    r2 = r2_score(y_test, y_pred)
    print(f'Mean Squared Error: {mse}')
    print(f'R^2 Score: {r2}')
    return history, model, mse, r2
```





# Plan B.2 - Skip sentiment analysis

What if I skipped the 0/1 or -1/1 step and embedded the headline text but otherwise did the same thing?

Keep more information in the input, plus more features, NN makes more sense...

Atrocious  $r^2$  values

	ticker	shift_days	mse	r2
0	TSLA	0	0.904176	-0.073068
1	TSLA	1	5.381370	-3.256683
2	TSLA	weekly	59.469277	-0.004637
3	AUR	weekly	38.998601	0.010772
4	GOOG	0	2.808688	-0.679575
5	GOOG	1	2.193855	-0.200059
6	GOOG	weekly	14.118953	-0.011923
7	GM	0	2.423245	-0.324621
8	GM	1	7.339129	-0.813074
9	GM	weekly	21.313594	-0.024720
10	NVDA	0	14.398336	-0.060567
11	NVDA	1	6.202351	-0.032551
12	NVDA	weekly	29.647777	-0.018149
13	APTV	0	4.369776	-0.951300
14	APTV	1	3.861646	-5.950820

# Plan B.3 - Random Forest

Models not returning anything useful

- Models are bad?
- Data is bad?
- No correlation?

Not having any luck with NN, try random forest model




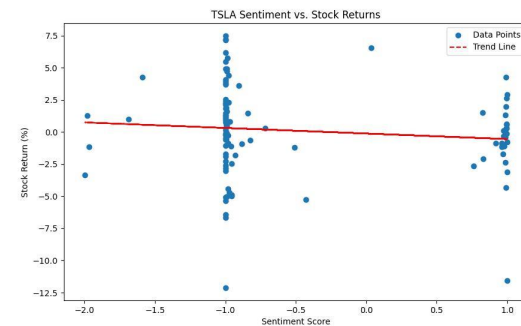
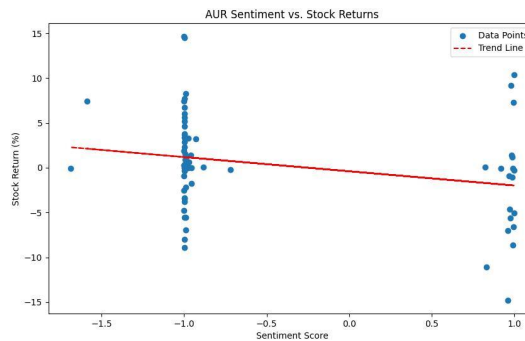


```
def train_random_forest(X_train, y_train, X_test, y_test):  
    model = RandomForestRegressor(n_estimators=100, random_state=777)  
    model.fit(X_train, y_train)  
    y_pred = model.predict(X_test)  
    mse = mean_squared_error(y_test, y_pred)  
    r2 = r2_score(y_test, y_pred)  
    print(f'Mean Squared Error: {mse}')  
    print(f'R^2 Score: {r2}')  
    return model, mse, r2
```

	ticker	shift_days	mse	r2
0	TSLA	0	0.871382	-0.034149
1	TSLA	1	5.434484	-3.298697
2	TSLA	weekly	62.870552	-0.062096
3	AUR	weekly	27.995865	0.289865
4	GOOG	0	2.096480	-0.253680
5	GOOG	1	1.789170	0.021307
6	GOOG	weekly	14.382935	-0.030843
7	GM	0	2.458829	-0.344072
8	GM	1	9.256998	-1.286869
9	GM	weekly	22.055687	-0.060398
10	NVDA	0	12.602407	0.071719
11	NVDA	1	6.509131	-0.083623
12	NVDA	weekly	31.669994	-0.087594
13	APTV	0	3.905757	-0.744095
14	APTV	1	3.478999	-5.262069
15	APTV	weekly	19.823030	-0.022911
16	TSPH	weekly	236.701643	-3.039999

# Hang on...

No Correlation !!

Forgot simple check




# Plan B.3 - Sentiment Lagged in Time LSTM

Because the no correlation graphs were on same day, I attempted to engineer features which were sentiment lagging by up to 7 days

Same types of results:

	ticker	mse	r2
0	TSLA	6.678150	-0.033874
1	AUR	24.779505	0.013853
2	GOOG	3.032027	-0.012493
3	GM	4.464870	-0.051762
4	NVDA	7.637391	0.023503
5	APTV	3.667235	-0.051548
6	TSPH	254.823822	NaN



# Plan C - "Blood in the Water"

Last ditch effort to see if the stock prices led articles sentiment instead.

Engineered lagging price pct\_change features by 2 weeks for 5 stocks

LSTM model

Accuracy: 0.6428571428571429

Confusion Matrix:

[[54 0]

[30 0]]

Classification Report:

	precision	recall	f1-score	support
0	0.64	1.00	0.78	54
1	0.00	0.00	0.00	30
accuracy			0.64	84
macro avg	0.32	0.50	0.39	84
weighted avg	0.41	0.64	0.50	84

	date	text					\	
18	2016-04-12	Facebook's Developer Conference Kicks Off On T...						
19	2016-04-13	Ford's Planned New Headquarters Borrow Some Si...						
20	2016-04-22	Despite Strong Earnings, G.M. Has Much to Prov...						
21	2016-04-27	Ford and Google Team Up to Support Driverless ...						
22	2016-05-03	Google to Get Fiat Chrysler Minivans for Self-...						
	sentiment_label	sentiment_score	sentiment	sentiment_binary	APTV		\	
18	POSITIVE	0.686108	1	1	-0.371802			
19	NEGATIVE	0.677211	-1	0	5.625432			
20	NEGATIVE	0.997127	-1	0	-0.355260			
21	NEGATIVE	0.871993	-1	0	-1.091558			
22	POSITIVE	0.997987	1	1	-0.732601			
	GM	GOOGL	NVDA	...	APTV_lag_5	APTV_lag_6	APTV_lag_7	\
18	0.609544	0.895004	-0.111485	...	2.290081	-2.219791	-0.943398	
19	3.601486	0.993043	2.511165	...	3.284344	2.290081	-2.219791	
20	-1.469686	-5.414102	-0.384508	...	2.793622	3.284344	2.290081	
21	-0.093193	-0.539034	2.275224	...	-3.881121	2.793622	3.284344	
22	-1.574803	-0.835655	-1.194445	...	-3.881121	-3.881121	2.793622	
	APTV_lag_8	APTV_lag_9	APTV_lag_10	APTV_lag_11	APTV_lag_12	\		
18	-2.808986	-1.040630	-1.173528	-1.945627	-5.598187			
19	-0.943398	-2.808986	-1.040630	-1.173528	-1.945627			
20	-2.219791	-0.943398	-2.808986	-1.040630	-1.173528			
21	2.290081	-2.219791	-0.943398	-2.808986	-1.040630			
22	3.284344	2.290081	-2.219791	-0.943398	-2.808986			
	APTV_lag_13	APTV_lag_14	\					
18	-2.657364	-2.774135						
19	-5.598187	-2.657364						
20	-1.945627	-5.598187						
21	-1.173528	-1.945627						
22	-1.040630	-1.173528						



# Lessons Learned

## LARGER DATA SETS

- Most of my problems really seem to have stemmed from datasets being too small. It's hard to do ML projects on a time crunch with "go and get it" data

## Clear Plan

- Genuinely went in with what I thought was a clear plan. But what exactly I wanted models to predict was not concrete and that led to some confusion on what models or parameters to select/try.

## Don't forget early data checks

- By the time I realized how futile my effort was, it was too late to switch horses. Could have saved myself a lot of time and compute.



# Loose ends

Would still like to see the downstream impact of headline sentiment on SOCIAL sentiment

- Didn't have time + Twitter 'pull' access API is \$100/mo

Would have liked to pull articles/headlines/text from more sources to build a much larger and more diverse corpus in the same time window